

# Exploration-Exploitation tradeoffs and multi-armed bandit problems

CS7032: AI & Agents for IET

November 24, 2015

# Learning and Feedback

- ▶ Consider the **ACO algorithm**: how does the system **learn**?
- ▶ **Contrast** that form of “learning” with, say, our system that learns to play **draughts**, or to a system that learns to filter out **spam mail**.

# Learning and Feedback

- ▶ Consider the **ACO algorithm**: how does the system **learn**?
- ▶ **Contrast** that form of “learning” with, say, our system that learns to play **draughts**, or to a system that learns to filter out **spam mail**.
- ▶ The RL literature often contrasts **instruction** and **evaluation**.
- ▶ **Evaluation** is a key component of Reinforcement learning systems:
  - ▶ Evaluative feedback is **local** (it indicates how good an action is) but not whether it is the best action possible
  - ▶ This creates a need for **exploration**

# Associative vs. non-associative settings

- ▶ In general, learning is both
  - ▶ **selectional**: i.e. actions are selected by trying different alternatives and comparing their effects,

# Associative vs. non-associative settings

- ▶ In general, learning is both
  - ▶ **selectional**: i.e. actions are selected by trying different alternatives and comparing their effects,
  - ▶ **associative**: i.e. the actions selected are associated to particular situations.

# Associative vs. non-associative settings


- ▶ In general, learning is both
  - ▶ **selectional**: i.e. actions are selected by trying different alternatives and comparing their effects,
  - ▶ **associative**: i.e. the actions selected are associated to particular situations.
- ▶ However, in order to study evaluative feedback in detail it is convenient to **simplify things** and consider the problem from a **non-associative** perspective.

# Learning from fruit machines

- ▶ The **n-armed bandit** setting:
  - ▶ Choice of  $n$  actions (which yield numerical rewards drawn from a stationary probability distribution)
    - ▶ each action selection called a **play**
  - ▶ Goal: **maximise** expected (long term) **total reward** or **return**.
    - ▶ Strategy: concentrate **plays** on the **best levers**.

# Learning from fruit machines

- ▶ The **n-armed bandit** setting:




- ▶ Choice of  $n$  actions (which yield numerical rewards drawn from a stationary probability distribution)
  - ▶ each action selection called a **play**
- ▶ Goal: **maximise** expected (long term) **total reward** or **return**.
  - ▶ Strategy: concentrate **plays** on the **best levers**.



# Learning from fruit machines


- ▶ The **n-armed bandit** setting:



- ▶ Choice of  $n$  actions (which yield numerical rewards drawn from a stationary probability distribution)
  - ▶ each action selection called a **play**
- ▶ Goal: **maximise** expected (long term) **total reward** or **return**.
  - ▶ Strategy: concentrate **plays** on the **best levers**.

# Learning from fruit machines


- ▶ The **n-armed bandit** setting:



- ▶ Choice of  $n$  actions (which yield numerical rewards drawn from a stationary probability distribution)
  - ▶ each action selection called a **play**
- ▶ Goal: **maximise** expected (long term) **total reward** or **return**.
  - ▶ Strategy: concentrate **plays** on the **best levers**.

# Learning from fruit machines


- ▶ The **n-armed bandit** setting:



- ▶ Choice of  $n$  actions (which yield numerical rewards drawn from a stationary probability distribution)
  - ▶ each action selection called a **play**
- ▶ Goal: **maximise** expected (long term) **total reward** or **return**.
  - ▶ Strategy: concentrate **plays** on the **best levers**.

# Learning from fruit machines

- ▶ The **n-armed bandit** setting:



- ▶ Choice of  $n$  actions (which yield numerical rewards drawn from a stationary probability distribution)
  - ▶ each action selection called a **play**
- ▶ Goal: **maximise** expected (long term) **total reward** or **return**.
  - ▶ Strategy: concentrate **plays** on the **best levers**.

# How to find the best plays

- ▶ Let's **distinguish** between:
  - ▶ **reward**, the immediate outcome of a play, and
  - ▶ **value**, the expected (mean) reward of a play
- ▶ But how do we **estimate** values?
- ▶ We could **keep a record of rewards**  $r_1^a, \dots, r_{k_a}^a$  for each chosen action  $a$  and estimate the value of choosing  $a$  at time  $t$  as

$$Q_t(a) = \frac{r_1^a + r_2^a + \dots + r_{k_a}^a}{k_a} \quad (1)$$

## But how do we choose an action?

- ▶ We could be **greedy** and **exploit** our knowledge of  $Q_t(a)$  to choose at any time  $t$  the play with the highest estimated value.
- ▶ But this strategy will tend to neglect the estimates of non-greedy actions (which might have produced greater total reward in the long run).
- ▶ One could, alternatively, **explore** the space of actions by choosing, from time to time, a non-greedy action

# Balancing exploitation and exploration

- ▶ The goal is to approximate the **true value**,  $Q(a)$  for each action
- ▶  $Q_t(a)$ , given by equation (1) will **converge** to  $Q(a)$  as  $k_a \rightarrow \infty$  (Law of Large Numbers)
- ▶ So balancing exploration and exploitation is necessary...
- ▶ but **finding the right balance can be tricky**; many approaches rely on strong assumptions about the underlying distributions
- ▶ We will present a simpler alternative...

# A simple strategy

- ▶ The **simplest strategy**: choose action  $a^*$  so that:

$$a^* = \arg \max_a Q_t(a) \quad (\text{greedy selection}) \quad (2)$$

- ▶ A better simple strategy:  $\epsilon$ -greedy methods:
  - ▶ With probability  $1 - \epsilon$ , exploit; i.e. **choose  $a^*$**  according with (2)
  - ▶ The rest of the time (probability  $\epsilon$ ) choose an action at **random**, uniformly
  - ▶ A suitably **small  $\epsilon$**  guarantees that all actions get **explored sooner or later** (and the **probability of selecting the optimal action** converges to greater than  $1 - \epsilon$ )



# Exploring exploration at different rates

- ▶ The “10-armed testbed” [Sutton and Barto, 1998]:
  - ▶ 2000 randomly generated  $n$ -armed bandit tasks with  $n = 10$ .
  - ▶ For each action,  $a$ , expected rewards  $Q(a)$  (the “true” expected values of choosing  $a$ ) are selected from a normal (Gaussian) probability distribution  $N(0, 1)$
  - ▶ Immediate rewards  $r_1^a, \dots, r_k^a$  are similarly selected from  $N(Q(a), 1)$



# Exploring exploration at different rates

- ▶ The “10-armed testbed” [Sutton and Barto, 1998]:
  - ▶ 2000 randomly generated  $n$ -armed bandit tasks with  $n = 10$ .
  - ▶ For each action,  $a$ , expected rewards  $Q(a)$  (the “true” expected values of choosing  $a$ ) are selected from a normal (Gaussian) probability distribution  $N(0, 1)$
  - ▶ Immediate rewards  $r_1^a, \dots, r_k^a$  are similarly selected from  $N(Q(a), 1)$

```
1 nArmedBanditTask <- function(n=10, s=2000)
2 {
3   qtrue <- rnorm(n)
4   r <- matrix(nrow=s, ncol=n)
5   for (i in 1:n){
6     r[, i] <- rnorm(s, mean=qtrue[i])
7   }
8   return(r)
9 }
```

# Some empirical results

- Averaged results of 2000 rounds of 1000 plays each:



# Softmax methods

- ▶ **Softmax action selection** methods grade action probabilities by estimated values.
- ▶ Commonly use **Gibbs, or Boltzmann, distribution**. I.e. choose action  $a$  on the  $t$ -th play with probability

$$\pi(a) = \frac{e^{Q_t(a)/\tau}}{\sum_{b \in A} e^{Q_t(b)/\tau}} \quad (3)$$

- ▶ where  $\tau$  is a positive parameter called the *temperature*, as in **simulated annealing** algorithms.

# Keeping track of the means

- Maintaining a record of means by recalculating them after each play can be a source of inefficiency in evaluating feedback. E.g.:

```
1  ## inefficient way of selecting actions
2  selectAction <- function(r, e=0) {
3    n <- dim(r)[2] # number of arms (cols in r)
4    if (sample(c(T,F),1,prob=c(e,1-e))) {
5      return(sample(1:n,1)) # explore if e > 0 and chance allows
6    }
7    else { # all Q_t(a_i), i \in [1,n]
8      Qt <- sapply(1:n, function(i) mean(r[,i], na.rm=T))
9      ties <- which(Qt == max(Qt))
10     ## if there are ties in Q_t(a); choose one at random
11     if (length(ties)>1)
12       return(sample(ties, 1))
13     else
14       return(ties)
15   }
16 }
```

## Incremental update

- ▶ Do we really need to store all rewards?

## Incremental update

- ▶ Do we really need to store all rewards? NO

## Incremental update

- ▶ Do we really need to store all rewards? NO
- ▶ Let  $Q_k$  be the mean of the first  $k$  rewards for an action,
- ▶ We can express the next mean value as

$$\begin{aligned} Q_{k+1} &= \frac{1}{k+1} \sum_{i=1}^{k+1} r_i \\ &= \dots \\ &= Q_k + \frac{1}{k+1} [r_{k+1} - Q_k] \end{aligned} \tag{4}$$



## Incremental update

- ▶ Do we really need to store all rewards? NO
- ▶ Let  $Q_k$  be the mean of the first  $k$  rewards for an action,
- ▶ We can express the next mean value as

$$\begin{aligned} Q_{k+1} &= \frac{1}{k+1} \sum_{i=1}^{k+1} r_i \\ &= \dots \\ &= Q_k + \frac{1}{k+1} [r_{k+1} - Q_k] \end{aligned} \tag{4}$$

- ▶ Equation (4) is part of family of formulae in which the **new estimate** ( $Q_{k+1}$ ) is the old estimate ( $Q_k$ ) adjusted by the **estimation error** ( $r_{k+1} - Q_k$ ) scaled by a **step parameter** ( $\frac{1}{k+1}$ , in this case)
- ▶ (Recall the LMS weight update step described in the introduction to machine learning).

# Non-stationary problems

- ▶ What happens if the “n-armed bandit” changes over time? How do we keep track of a changing mean?
- ▶ An approach: weight recent rewards more heavily than past ones.
- ▶ So, our update formula (4) could be modified to

$$Q_{k+1} = Q_k + \alpha[r_{k+1} - Q_k] \quad (5)$$

where  $0 < \alpha \leq 1$  is a constant

- ▶ This gives us  $Q_k$  as an **exponential, recency-weighted average** of past rewards and the initial estimate

# Recency weighted estimates

- ▶ Given (5),  $Q_k$  can be rewritten as:

$$\begin{aligned}Q_k &= Q_{k-1} + \alpha[r_k - Q_{k-1}] \\&= \alpha r_k + (1 - \alpha)Q_{k-1} \\&= \dots \\&= (1 - \alpha)^k Q_0 + \sum_{i=1}^k \alpha(1 - \alpha)^{k-i} r_i\end{aligned}\tag{6}$$

- ▶ Note that  $(1 - \alpha)^k + \sum_{i=1}^k \alpha(1 - \alpha)^{k-i} = 1$ . I.e. weights sum to 1.
- ▶ The weight  $\alpha(1 - \alpha)^{k-i}$  decreases exponentially with the number of intervening rewards

# Picking initial values

- ▶ The above methods are **biased** by  $Q_0$ 
  - ▶ For sample-average, bias disappears once all actions have been selected
  - ▶ For recency-weighted methods, bias is **permanent** (but decreases over time)
- ▶ One can supply **prior knowledge** by picking the right values for  $Q_0$
- ▶ One can also choose **optimistic initial values** to encourage exploration

# Picking initial values

- ▶ The above methods are **biased** by  $Q_0$ 
  - ▶ For sample-average, bias disappears once all actions have been selected
  - ▶ For recency-weighted methods, bias is **permanent** (but decreases over time)
- ▶ One can supply **prior knowledge** by picking the right values for  $Q_0$
- ▶ One can also choose **optimistic initial values** to encourage exploration (**Why?**)


# The effects of optimism

# The effects of optimism

- ▶ With  $Q_0$ 's set initially high (for each action) $q$ , the learner will be **disappointed** by actual rewards, and sample all actions many times before converging.

# The effects of optimism

- ▶ With  $Q_0$ 's set initially high (for each action) $q$ , the learner will be **disappointed** by actual rewards, and sample all actions many times before converging.
- ▶ E.g.: Performance of **Optimistic** ( $Q_0 = 5, \forall a$ ) vs **Realistic** ( $Q_0 = 0, \forall a$ ) strategies for the 10-armed testbed



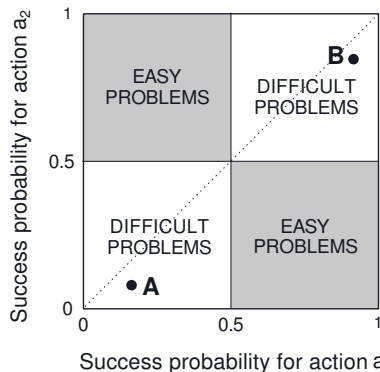


# Evaluation vs Instruction

- ▶ RL searches the **action space** while SL searches the **parameter space**.
- ▶ **Binary**, 2-armed bandits:
  - ▶ two actions:  $a_1$  and  $a_2$
  - ▶ two rewards: **success** and **failure** (as opposed to numeric rewards).
- ▶ **SL**: select the **action** that returns **success most often**.

# Evaluation vs Instruction

- ▶ RL searches the **action space** while SL searches the **parameter space**.
- ▶ **Binary**, 2-armed bandits:
  - ▶ two actions:  $a_1$  and  $a_2$
  - ▶ two rewards: **success** and **failure** (as opposed to numeric rewards).
- ▶ **SL**: select the **action** that returns **success most often**.
- ▶ Stochastic case (a **problem** for SL?):



## Two stochastic SL schemes (learning automata)

- ▶  $L_{r-p}$ , **Linear reward-penalty**: choose the action as in the naive SL case, and keep track of the successes by updating an estimate of the probability of choosing it in future.
- ▶ Choose  $a \in A$  with **probability**  $\pi_t(a)$
- ▶ Update the probability **of the chosen action** as follows:

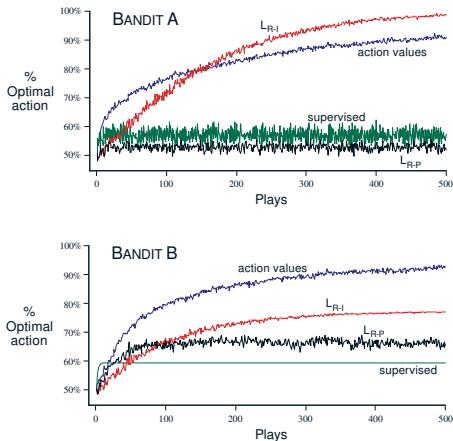
$$\pi_{t+1}(a) = \begin{cases} \pi_t(a) + \alpha(1 - \pi_t(a)) & \text{if success} \\ \pi_t(a)(1 - \alpha) & \text{otherwise} \end{cases} \quad (7)$$

- ▶ All **remaining actions**  $a_i \neq a$  are adjusted proportionally:

$$\pi_{t+1}(a_i) = \begin{cases} \pi_t(a_i)(1 - \alpha) & \text{if a succeeds} \\ \frac{\alpha}{|A|-1} + \pi_t(a_i)(1 - \alpha) & \text{otherwise} \end{cases} \quad (8)$$

- ▶  $L_{r-i}$ , **Linear reward-inaction**: like  $L_{r-p}$  but update probabilities only in case of success.

# Performance: instruction vs evaluation



- See [Narendra and Thathachar, 1974], for a survey of learning automata methods and results.

# Reinforcement comparison

- Instead of estimates of values for each action, keep an estimate of overall reward level  $\bar{r}_t$ :

$$\bar{r}_{t+1} = \bar{r}_t + \alpha[r_{t+1} - r_t] \quad (9)$$

- and **action preferences**  $p_t(a)$  w.r.t. reward estimates:


$$p_{t+1}(a) = p_t(a) + \beta[r_t - \bar{r}_t] \quad (10)$$

- The **softmax function** can then be used as a PMF for action selection:

$$\pi_t(a) = \frac{e^{p_t(a)}}{\sum_{b \in A} e^{p_t(b)}} \quad (11)$$

- (  $0 < \alpha \leq 1$  and  $\beta > 0$  are step-size parameters.)

# How does reinforcement comparison compare?



- Reinforcement comparison ( $\alpha = 0.1$ ) vs action-value on the 10-armed testbed.

# Pursuit methods


- ▶ maintain both action-value estimates and action preferences,
- ▶ preferences continually "pursue" greedy actions.
- ▶ E.g.: for greedy action  $a^* = \arg \max_a Q_{t+1}(a)$ , increase its selection probability:

$$\pi_{t+1}(a^*) = \pi_t(a^*) + \beta[1 - \pi_t(a^*)] \quad (12)$$

- ▶ while decreasing probabilities for all remaining actions  $a \neq a^*$

$$\pi_{t+1}(a) = \pi_t(a) + \beta[0 - \pi_t(a)] \quad (13)$$

# Performance of pursuit methods



- Pursuit ( $\alpha = 1/k$  for  $Q_t$  update,  $\pi_0(a) = 1/n$  and  $\beta = 0.01$ ) versus reinforcement comparison ( $\alpha = 0.1$ ) vs action-value on the 10-armed testbed.



## Further topics

- ▶ Associative search:
  - ▶ suppose we have many different bandit tasks and the learner is presented with a different one at each play.
  - ▶ Suppose each time the learner is given a clue as to the which task it is facing...

# Further topics

- ▶ Associative search:
  - ▶ suppose we have many different bandit tasks and the learner is presented with a different one at each play.
  - ▶ Suppose each time the learner is given a clue as to the which task it is facing...
  - ▶ In such cases, the best strategy is a combination of *search* for the best actions and *association* of actions to situations

## Further topics

- ▶ Associative search:
  - ▶ suppose we have many different bandit tasks and the learner is presented with a different one at each play.
  - ▶ Suppose each time the learner is given a clue as to the which task it is facing...
  - ▶ In such cases, the best strategy is a combination of *search* for the best actions and *association* of actions to situations
- ▶ Exact algorithms for computation of *Bayes optimal* way to balance exploration and exploitation exist [Bellman, 1956], but are intractable.

# References



Bellman, R. (1956).

A problem in the sequential design of experiments.  
*Sankhya*, 16:221–229.



Narendra, K. and Thathachar, M. (1974).

Learning automata - a survey.  
*IEEE Transactions on Systems, Man and Cybernetics*, SMC-4(4):323–334.



Sutton, R. S. and Barto, A. G. (1998).

*Reinforcement Learning: An Introduction*.  
MIT Press, Cambridge, MA.