

# Combining Dynamic programming and approximation architectures

AI & Agents for IET


Lecturer: Liliana Mamani Sanchez

<http://www.scss.tcd.ie/~mamanisl/teaching/cs7032/>

November 30, 2015

# Combining TD and function approximation

- ▶ Basic idea: use supervised learning to provide an approximation of the value function for TD learning
- ▶ The approximation architecture is should generalise over (possibly unseen) states



- ▶ In a sense, it groups states into equivalence classes (wrt value)

# Why use approximation architectures

- ▶ To cope with the **curse of dimensionality**
- ▶ by generalising over **states**
  - ▶ Note that the algorithms we have seen so far (DP, TD, Sarsa, Q-learning) all use tables to store states (or state-action tuples)
  - ▶ This works well if the number of states is relatively small
  - ▶ But it doesn't scale up very well
- ▶ (We have already seen **examples** of approximation architectures: the draughts player, the examples in the neural nets lecture.)

# Gradient descent methods

- ▶ The **LMS algorithm** use for draughts illustrates a gradient descent method
  - ▶ (to approximate a linear function)
- ▶ **Goal**: to learn the **parameter** vector

$$\vec{\theta}_t = (\theta_t(1), \theta_t(2), \theta_t(3), \dots, \theta_t(m)) \quad (1)$$

by adjusting them at each iteration towards **reducing the error**:

$$\vec{\theta}_{t+1} = \vec{\theta}_t - \frac{1}{2} \alpha \nabla_{\vec{\theta}_t} (V^\pi(s_t) - V_t(s_t))^2 \quad (2)$$

$$= \vec{\theta}_t + (V^\pi(s_t) - V_t(s_t)) \alpha \nabla_{\vec{\theta}_t} V_t(s_t) \quad (3)$$

where  $V_t$  is a smooth, differentiable function of  $\vec{\theta}_t$ .

## Backward view and update rule

- ▶ The problem with (2) is that the target value ( $V^\pi$ ) is typically not available.
- ▶ Different methods replace their estimates for this value function:
  - ▶ So Monte Carlo, for instance, would use the return  $R_t$
  - ▶ And the  $TD(\lambda)$  method uses  $R_t^\lambda$ :

$$\vec{\theta}_{t+1} = \vec{\theta}_t + \alpha(R_t^\lambda - V_t(s_t))\nabla_{\vec{\theta}_t} V_t(s_t) \quad (4)$$

- ▶ The **backward view** is given by:

$$\vec{\theta}_{t+1} = \vec{\theta}_t + \alpha\delta_t\vec{e}_t \quad (5)$$

where  $\vec{e}_t$  is a vector of eligibility traces (one for each component of  $\vec{\theta}_t$ ), updated by

$$\vec{e}_t = \gamma\lambda\vec{e}_{t-1} + \nabla_{\vec{\theta}_t} V_t(s_t) \quad (6)$$

# Value estimation with approximation

## Algorithm 1: On-line gradient descent TD( $\lambda$ )

```
1   Initialise  $\vec{\theta}$  arbitrarily
2        $\vec{e} \leftarrow 0$ 
3        $s \leftarrow$  initial state of episode
4   repeat (for each step of episode)
5       choose  $a$  according to  $\pi$ 
6       perform  $a$ , observe  $r, s'$ 
7        $\delta \leftarrow r + \gamma V(s') - V(s)$ 
8        $\vec{e} \leftarrow \gamma\lambda\vec{e} + \nabla_{\vec{\theta}} V(s)$ 
9        $\vec{\theta} \leftarrow \vec{\theta} + \alpha\delta\vec{e}$ 
10       $s \leftarrow s'$ 
11  until  $s$  is terminal state
```

► Methods commonly used to compute the **gradients**  $\nabla_{\vec{\theta}} V(s)$ :

- error **back-propagation** (multilayer NNs), or by
- **linear approximators** (for value functions of the form  $V_t(s) = (\vec{\theta}_t)^T \vec{f} = \sum_{i=1}^n \theta_t(i) f(i)$ . (where  $(\vec{\theta}_t)^T$  denotes the transpose of  $\vec{\theta}_t$ )

# Control with approximation

- ▶ The general (**forward view**) update rule for action-value prediction (by gradient descent) can be written:

$$\vec{\theta}_{t+1} = \vec{\theta}_t + \alpha(R_t^\lambda - Q_t(s_t, a_t))\nabla_{\vec{\theta}_t} Q_t(s_t, a_t) \quad (7)$$

(recal that  $V_t$  is determined by  $\vec{\theta}_t$ )

- ▶ So the **backward view** can be expressed as before:

$$\vec{\theta}_{t+1} = \vec{\theta}_t + \alpha\delta_t\vec{e}_t \quad (8)$$

where

$$\vec{e}_t = \gamma\lambda\vec{e}_{t-1} + \nabla_{\vec{\theta}_t} Q_t(s_t, a_t) \quad (9)$$

# An algorithm: gradient descent Q-learning


## Algorithm 2: Linear Gradient Descent $Q(\lambda)$

```
1   Initialise  $\theta$  arbitrarily
2   for each episode
3      $\vec{e} \leftarrow \vec{0}$ ; initialise  $s, a$ 
4      $\mathcal{F}_a \leftarrow$  set of features in  $s, a$ 
5     repeat (for each step of episode)
6       for all  $i \in \mathcal{F}_a$ :  $e(i) \leftarrow e(i) + 1$ 
7       perform  $a$ , observe  $r, s$ 
8        $\delta \leftarrow r - \sum_{i \in \mathcal{F}_a} \theta(i)$ 
9       for all  $a \in A$ 
10         $\mathcal{F}_a \leftarrow$  set of features in  $s, a$ 
11         $Q_a \leftarrow \sum_{i \in \mathcal{F}_a} \theta(i)$ 
12         $\delta \leftarrow \delta + \gamma \max_a Q_a$ 
13         $\vec{\theta} \leftarrow \vec{\theta} + \alpha \delta \vec{e}$ 
14        with probability  $1 - \epsilon$ 
15          for all  $a \in A$ 
16             $Q_a \leftarrow \sum_{i \in \mathcal{F}_a} \theta(i)$ 
17             $a \leftarrow \arg \max_a Q_a$ 
18             $\vec{e} \leftarrow \gamma \lambda \vec{e}$ 
19        else
20           $a \leftarrow$  a random action
21           $\vec{e} \leftarrow \vec{0}$ 
22   until  $s$  is terminal state
```



# A Case Study: TD-Gammon

- ▶ 15 white and 15 black pieces on a board of 24 locations, called points.
- ▶ Player rolls 2 dice and can move 2 pieces (or same piece twice)
- ▶ Goal is to move pieces to last quadrant (for white that's 19-24) and then off the board
- ▶ A player can “hit” any opposing single piece placed on a point, causing that piece to be moved to the “bar”
- ▶ Two pieces on a point block that point for the opponent
- ▶ + a number of other complications

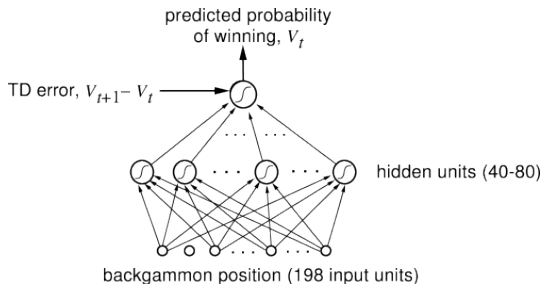


# Game complexity

- ▶ 30 pieces, 26 locations
- ▶ Large number of **actions** possible from a given state (up to 20)
- ▶ Very large number of possible **states** ( $10^{20}$ )
- ▶ Branching factor of about 400 (so difficult to apply heuristics)
- ▶ **Stochastic** environment (next state depends on the opponent's move) but **fully observable**


# TD-Gammon's solution

- ▶  $V_t(s)$  meant to estimate the **probability of winning** from any state  $s$
- ▶ **Rewards**: 0 for all stages, except those on which the game is won
- ▶ Learning: non-linear form of  $TD(\lambda)$ 
  - ▶ like the Algorithm presented above, using a multilayer neural network to compute the gradients



# State representation in TD-Gammon

- ▶ Representation involved little domain knowledge
- ▶ 198 input features:



- ▶ For each point on the backgammon board, **four units** indicated the number of white pieces on the point (see [Tesauro, 1994] for a detailed description of the encoding used)
- ▶  $(4 \text{ (white)} + 4 \text{ (black)}) \times 24 \text{ points} = 192 \text{ units}$
- ▶ 2 units encoded the number of white and black pieces on the bar
- ▶ 2 units encoded the number of black and white pieces already successfully removed from the board
- ▶ 2 units indicated in a binary fashion whether it was white's or black's turn to move.

# TD-Gammon learning

- ▶ Given state (position) representation, the network computed its estimate in the way described in lecture 10.
  - ▶ Output of hidden unit  $j$  given by a sigmoid function of the weighted sum of inputs  $i$

$$h(j) = \sigma\left(\sum_i w_{ij} f(i)\right) \quad (10)$$

- ▶ Computation from hidden to output units is analogous to this
- ▶ TD-Gammon employed  $TD(\lambda)$  where the eligibility trace updates (equation (9)),

$$\vec{e}_t = \gamma \lambda \vec{e}_{t-1} + \nabla_{\vec{\theta}_t} V_t(s_t)$$

were computed by the back-propagation procedure

- ▶ TD-Gammon set  $\gamma = 1$  and rewards to zero, except on winning, so TD error is usually  $V_t(s_{t+1}) - V_t(s_t)$

# TD-Gammon training

- ▶ Training data obtained by playing against itself
- ▶ Each game was treated as an episode
- ▶ Non-linear TD applied incrementally (i.e. after each move)
- ▶ Some results (according to [Sutton and Barto, 1998])

Program	Hidden Units	Training Games	Opponents	Results
TD-Gam 0.0	40	300,000	other programs	tied for best
TD-Gam 1.0	80	300,000	Robertie, Magriel, ...	-13 pts / 51 games
TD-Gam 2.0	40	800,000	various Grandmasters	-7 pts / 38 games
TD-Gam 2.1	80	1,500,000	Robertie	-1 pt / 40 games
TD-Gam 3.0	80	1,500,000	Kazaros	+6pts / 20 games

# References

Notes based on [Sutton and Barto, 1998, ch 8, 9]. Further details on TD-Gammon can be found in Tesauro's papers [Tesauro, 1994]. Other interesting case studies can be found in [Sutton and Barto, 1998, ch 10] and [Bertsekas and Tsitsiklis, 1996].



Bertsekas, D. P. and Tsitsiklis, J. N. (1996).

*Neuro-Dynamic Programming.*

Athena Scientific, Belmont.



Sutton, R. S. and Barto, A. G. (1998).

*Reinforcement Learning: An Introduction.*

MIT Press, Cambridge, MA.



Tesauro, G. (1994).

TD-gammon, a self-teaching backgammon program, achieves master-level play.

*Neural Computation*, 6:215–219.