

# Temporal difference learning

AI & Agents for IET

Lecturer: Liliana Mamani Sanchez

<http://www.scss.tcd.ie/~mamanisl/teaching/cs7032/>

November 30, 2015

## Recall background & assumptions

- ▶ Environment is a **finite MDP** (i.e.  $A$  and  $S$  are finite).
- ▶ MDP's dynamics defined by **transition probabilities**:

$$\mathcal{P}_{ss'}^a = P(s_{t+1} = s' | s_t = s, a_t = a)$$

- ▶ and expected **immediate rewards**,

$$\mathcal{R}_{ss'}^a = E\{r_{t+1} | s_t = s, a_t = a, s_{t+1} = s'\}$$

- ▶ Goal: to **search for good policies**  $\pi$
- ▶ **DP Strategy**: use **value functions** to structure search:

$$V^*(s) = \max_a \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma V^*(s')] \quad \text{or}$$

$$Q^*(s, a) = \sum_{s'} \mathcal{P}_{ss'}^a [\mathcal{R}_{ss'}^a + \gamma \max_{a'} Q^*(s', a')]$$

- ▶ **MC strategy**: **expand each episode** and keep averages of returns per state.

# Temporal Difference (TD) Learning

- ▶ TD is a combination of ideas from Monte Carlo methods and DP methods:
  - ▶ TD can learn directly from raw experience without a model of the environment's dynamics, like MC.
  - ▶ TD methods update estimates based in part on other learned estimates, without waiting for a final outcome (i.e. they “bootstrap”), like DP
- ▶ Some widely used TD methods:  $TD(0)$ ,  $TD(\lambda)$ , Sarsa, Q-learning, Actor-critic, R-Learning

# Prediction & Control

- ▶ Recall **value function estimation** (prediction) for DP and MC:
  - ▶ Updates **for DP**:
  - ▶ Update rule **for MC**:

# Prediction & Control

- ▶ Recall **value function estimation** (prediction) for DP and MC:
  - ▶ Updates **for DP**:

$$\begin{aligned}V^{\pi}(s) &= E_{\pi}\{R_t | s_t = s\} \\ &= E_{\pi}\{r_{t+1} + \gamma V^{\pi}(s_{t+1}) | s_t = s\}\end{aligned}$$

- ▶ Update rule **for MC**:

# Prediction & Control

- ▶ Recall **value function estimation** (prediction) for DP and MC:
  - ▶ Updates **for DP**:

$$\begin{aligned} V^\pi(s) &= E_\pi\{R_t | s_t = s\} \\ &= E_\pi\{r_{t+1} + \gamma V^\pi(s_{t+1}) | s_t = s\} \end{aligned}$$

- ▶ Update rule **for MC**:

$$V(s_t) \leftarrow V(s_t) + \alpha[R_t - \gamma V(s_t)]$$

# Prediction & Control

- ▶ Recall **value function estimation** (prediction) for DP and MC:
  - ▶ Updates **for DP**:

$$\begin{aligned}V^\pi(s) &= E_\pi\{R_t | s_t = s\} \\ &= E_\pi\{r_{t+1} + \gamma V^\pi(s_{t+1}) | s_t = s\}\end{aligned}$$

- ▶ Update rule **for MC**:

$$V(s_t) \leftarrow V(s_t) + \alpha[R_t - \gamma V(s_t)]$$

(Q: why are the value functions above merely **estimates**?)

# Prediction in TD(0)

- Prediction (update of **estimates of  $V$** ) for the TD(0) method is done as follows:

$$V(s_t) \leftarrow V(s_t) + \alpha[r_{t+1} + \gamma V(s_{t+1}) - V(s_t)] \quad (1)$$

- No need to **wait** until the end of the episode (as in MC) to update  $V$
- No need to **sweep** across the possible successor states (as in DP).



# Prediction in TD(0)

- Prediction (update of **estimates of  $V$** ) for the TD(0) method is done as follows:

$$V(s_t) \leftarrow V(s_t) + \alpha \left[ \underbrace{r_{t+1} + \gamma V(s_{t+1})}_{\text{1-step estimate of } R_t} - V(s_t) \right] \quad (1)$$

- No need to **wait** until the end of the episode (as in MC) to update  $V$
- No need to **sweep** across the possible successor states (as in DP).

# Tabular TD(0) value function estimation

- Use **sample backup** (as in MC) **rather than full backup** (as in DP):


## Algorithm 1: TD(0)

---

```
1  TabularTD0( $\pi$ )
2    Initialisation ,  $\forall s \in S$ :
3     $V(s) \leftarrow$  arbitrary;
4     $\alpha \leftarrow$  learning constant
5     $\gamma \leftarrow$  discount factor
6
7    For each episode  $E$  until stop-criterion met
8      Initialise  $s$ 
9      For each step of  $E$ 
10          $a \leftarrow \pi(s)$ 
11         Take action  $a$ ;
12         Observe reward  $r$  and next state  $s'$ 
13          $V(s) \leftarrow V(s) + \alpha[r + \gamma V(s') - V(s)]$ 
14          $s \leftarrow s'$ 
15     until  $s$  is a terminal state
```


---

## Example: random walk estimate

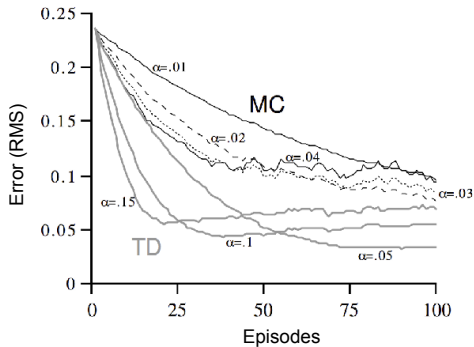


- ▶  $\mathcal{P}_{ss'}^a = .5$  for all non-terminal states
- ▶  $\mathcal{R}_{ss'}^a = 0$  except for terminal state on the right which is 1.
- ▶ Comparison between MC and TD(0) (over 100 episodes)

## Example: random walk estimate



- ▶  $\mathcal{P}_{ss'}^a = .5$  for all non-terminal states
- ▶  $\mathcal{R}_{ss'}^a = 0$  except for terminal state on the right which is 1.
- ▶ Comparison between MC and TD(0) (over 100 episodes)




# Focusing on action-value function estimates

- ▶ The **control** (algorithm) part.
- ▶ **Exploration** Vs. **exploitation** trade-off
- ▶ **On-policy** and **off-policy** methods
- ▶ Several variants, e.g.:
  - ▶ **Q-learning**: off-policy method [Watkins and Dayan, 1992]
  - ▶ Modified Q-learning (aka **Sarsa**): on-policy method

# Sarsa

- Transitions from non-terminal states **update**  $Q$  as follows:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)] \quad (2)$$



(for **terminal states**  $s_{t+1}$ , assign  $Q(s_{t+1}, a_{t+1}) \leftarrow 0$ )

## Algorithm 2: Sarsa

```
1  Initialise  $Q(s,a)$  arbitrarily
2  for each episode
3    initialise  $s$ 
4    choose  $a$  from  $s$  using  $\pi$  derived from  $Q$  /* e.g.  $\epsilon$ -greedy */
5    repeat (for each step of episode)
6      perform  $a$ , observe  $r, s'$ 
7      choose  $a'$  from  $s'$  using  $\pi$  derived from  $Q$  /* e.g.  $\epsilon$ -greedy */
8       $Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma Q(s',a') - Q(s,a)]$ 
9       $s \leftarrow s'$ 
10      $a \leftarrow a'$ 
11  until  $s$  is terminal state
```

# Q-learning

- An **off-policy method**: approximate  $Q^*$  independently of the policy being followed:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)] \quad (3)$$

## Algorithm 3: Q-Learning



---

```
1  Initialise  $Q(s, a)$  arbitrarily
2  for each episode
3      initialise  $s$ 
4      repeat (for each step of episode)
5          choose  $a$  from  $s$  using  $\pi$  derived from  $Q$  /* e.g.  $\epsilon$ -greedy */
6          perform  $a$ , observe  $r, s'$ 
7           $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ 
8           $s \leftarrow s'$ 
9  until  $s$  is terminal state
```

---

# A simple comparison

- ▶ A comparison between **Q-Learning** and **SARSA** on the “cliff walking” problem [Sutton and Barto, 1998]



- ▶ Why does Q-learning find the optimal path? Why is its average reward per episode worse than SARSA's?




# Convergence theorems

- Convergence, in the following sense:

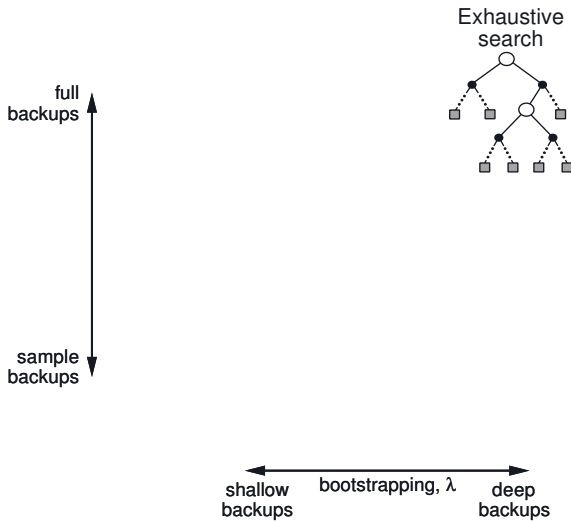
$Q$  converges in the limit to the optimal  $Q^*$  function, provided the system can be modelled as a deterministic Markov process,  $r$  is bounded and  $\pi$  visit every action-state pair infinitely often.

- Has been proved for the above described TD methods: Q-Learning [Watkins and Dayan, 1992], Sarsa and TD(0) [Sutton and Barto, 1998]. General proofs based on stochastic approximation theory can be found in [Bertsekas and Tsitsiklis, 1996].


# Summary of methods




# Summary of methods




# Summary of methods



# Summary of methods

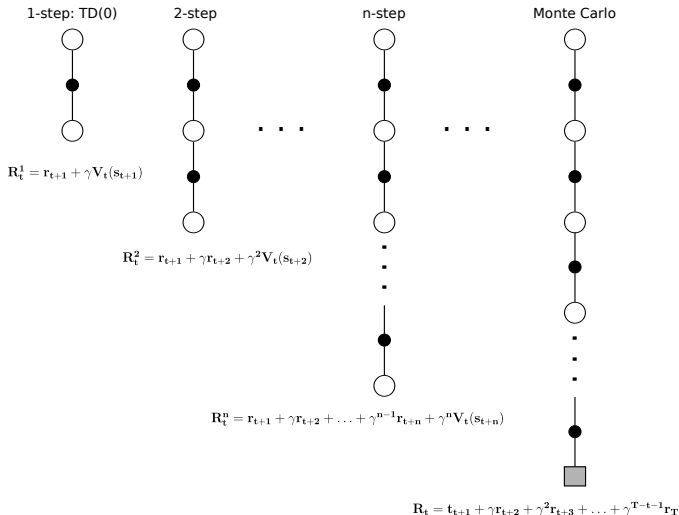


# Summary of methods



# TD( $\lambda$ )

- Basic Idea: if with TD(0) we backed up our estimates based on **one step ahead**, why not generalise this to include **n-steps** ahead?



# Averaging backups

- ▶ Consider, **for instance**, averaging 2- and 4-step backups:

$$R_t^\mu = 0.5R_t^2 + 0.5R_t^4$$

- ▶ TD( $\lambda$ ) is a method for averaging **all** n-step backups.
  - ▶ Weight by  $\lambda^{n-1}$  ( $0 \leq \lambda \leq 1$ ):

$$R_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} R_t^n$$


- ▶ **Backup** using  $\lambda$ -return:

$$\Delta V_t(s_t) = \alpha [R_t^\lambda - V_t(s_t)]$$

- ▶ [Sutton and Barto, 1998] call this the **Forward View of TD( $\lambda$ )**



# TD( $\lambda$ ) backup structure



$$\lim_{n \rightarrow \infty} \sum_{i=1}^n (1 - \lambda) \lambda^{i-1} = 1$$

Backup value ( $\lambda$ -return):

$$R_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} R_t^n$$

or

$$R_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} R_t^n + \lambda^{T-t-1} R_t$$

# Implementing TD( $\lambda$ )

- ▶ We need a way to accumulate the effect of the **trace-decay parameter**  $\lambda$
- ▶ Eligibility traces:

$$e_t(s) = \begin{cases} \gamma\lambda e_{t-1}(s) + 1 & \text{if } s = s_t \\ \gamma\lambda e_{t-1}(s) & \text{otherwise.} \end{cases} \quad (4)$$

- ▶ TD error for state-value prediction is:

$$\delta_t = r_{t+1} + \gamma V_t(s_{t+1}) - V_t(s_t) \quad (5)$$

- ▶ Sutton and Barto call this the **Backward View of TD( $\lambda$ )**

# Equivalence

- ▶ It can be shown [Sutton and Barto, 1998] that:

The Forward View of  $TD(\lambda)$  and the Backward View of  $TD(\lambda)$  are equivalent.

# Tabular TD( $\lambda$ )

## Algorithm 4: On-line TD( $\lambda$ )

---

```
1  Initialise  $V(s)$  arbitrarily
2       $e(s) \leftarrow 0$  for all  $s \in S$ 
3  for each episode
4      initialise  $s$ 
5      repeat (for each step of episode)
6          choose  $a$  according to  $\pi$ 
7          perform  $a$ , observe  $r, s'$ 
8           $\delta \leftarrow r + \gamma V(s') - V(s)$ 
9           $e(s) \leftarrow e(s) + 1$ 
10         for all  $s$ 
11              $V(s) \leftarrow V(s) + \alpha \delta e(s)$ 
12              $e(s) \leftarrow \gamma \lambda e(s)$ 
13          $s \leftarrow s'$ 
14     until  $s$  is terminal state
```

---

- Similar algorithms can be implemented for **control** (Sarsa, Q-learning), using eligibility traces.

# Control algorithms: SARSA( $\lambda$ )

## Algorithm 5: SARSA( $\lambda$ )

---

```
1   Initialise  $Q(s,a)$  arbitrarily
2        $e(s,a) \leftarrow 0$  for all  $s \in S$  and  $a \in A$ 
3   for each episode
4       initialise  $s,a$ 
5       repeat (for each step of episode)
6           perform  $a$ , observe  $r,s'$ 
7           choose  $a',s'$  according to  $Q$  (e.g.  $\epsilon$ -greedy)
8            $\delta \leftarrow r + \gamma Q(s',a') - Q(s,a)$ 
9            $e(s,a) \leftarrow e(s,a) + 1$ 
10          for all  $s$ 
11               $Q(s,a) \leftarrow Q(s,a) + \alpha \delta e(s,a)$ 
12               $e(s) \leftarrow \gamma \lambda e(s)$ 
13           $s \leftarrow s'$   $a \leftarrow a'$ 
14   until  $s$  is terminal state
```

---

# References

Notes based on [Sutton and Barto, 1998, ch 6].



Bertsekas, D. P. and Tsitsiklis, J. N. (1996).

*Neuro-Dynamic Programming.*

Athena Scientific, Belmont.



Sutton, R. S. and Barto, A. G. (1998).

*Reinforcement Learning: An Introduction.*

MIT Press, Cambridge, MA.



Watkins, C. J. C. H. and Dayan, P. (1992).

Technical note Q-learning.

*Machine Learning*, 8:279–292.