

# Dynamic Programming and Reinforcement Learning

## Lecture 4: Dimension Reduction in DP and RL

Mengdi Wang

Operations Research and Financial Engineering  
Princeton University

July 25-29, 2016

# Today

- ① Review: Infinite-Horizon DP
- ② Dimension Reduction in RL
  - Approximation in Value Space
  - Approximation in Policy Space
  - State Aggregation
- ③ On-Policy Learning
  - Direct Projection
  - Bellman Error Minimization
  - Projected Bellman Equation Method
  - From On-Policy to Off-Policy

# Infinite-Horizon DP: Theory

- Infinite horizon cost function expressions [with  $J_0(x) \equiv 0$ ]

$$J_\pi(x) = \lim_{N \rightarrow \infty} (T_{\mu_0} T_{\mu_1} \cdots T_{\mu_N} J_0)(x), \quad J_\mu(x) = \lim_{N \rightarrow \infty} (T_\mu^N J_0)(x)$$

- Bellman's equation:

$$J^* = T J^*, \quad J_\mu = T_\mu J_\mu$$

- Optimality condition:

$$\mu: \text{optimal} \quad \Longleftrightarrow \quad T_\mu J^* = T J^*$$

# Infinite-Horizon DP: Algorithms

- Value iteration: For any (bounded)  $J$

$$J^*(x) = \lim_{k \rightarrow \infty} (T^k J)(x), \quad \forall x$$

- Policy iteration: Given  $\mu^k$ ,
  - Policy evaluation: Find  $J_{\mu^k}$  by solving

$$J_{\mu^k} = T_{\mu^k} J_{\mu^k}$$

- Policy improvement : Find  $\mu^{k+1}$  such that

$$T_{\mu^{k+1}} J_{\mu^k} = T J_{\mu^k}$$

# Practical Difficulties of DP

## The curse of dimensionality

- Exponential growth of the computational and storage requirements as the number of state variables and control variables increases
- Quick explosion of the number of states in combinatorial problems

## The curse of modeling

- Sometimes a simulator of the system is easier to construct than a model
- There may be real-time solution constraints
- A family of problems may be addressed. The data of the problem to be solved is given with little advance notice
- The problem data may change as the system is controlled – need for on-line replanning

All of the above are motivations for approximation and simulation

# General Orientation to ADP

ADP (late 80s - present) is a breakthrough methodology that **allows the application of DP to problems with many or infinite number of states** .

Other names for ADP are:

- **“reinforcement learning”** (RL).
- **“neuro-dynamic programming”** (NDP).
- **“adaptive dynamic programming”** (ADP).

We will mainly adopt an  $n$ -state discounted model (the easiest case - but think of HUGE  $n$ ).

- Extensions to other DP models (continuous space, continuous-time, not discounted) are possible (but more quirky). We will set aside for later.

There are many approaches:

- Problem approximation
- Simulation-based approaches (we will focus on these)

Simulation-based methods are of three types:

- Rollout (we will not discuss further)
- Approximation in value space
- Approximation in policy space

- 1 Review: Infinite-Horizon DP
- 2 Dimension Reduction in RL
  - Approximation in Value Space
  - Approximation in Policy Space
  - State Aggregation
- 3 On-Policy Learning
  - Direct Projection
  - Bellman Error Minimization
  - Projected Bellman Equation Method
  - From On-Policy to Off-Policy

## Approximation in value space

- Approximate  $J^*$  or  $J_\mu$  from a parametric class  $\tilde{J}(i; r)$  where  $i$  is the current state and  $r = (r_1, \dots, r_m)$  is a vector of “tunable” scalar weights
- Use  $\tilde{J}$  in place of  $J^*$  or  $J_\mu$  in various algorithms and computations
- Role of  $r$  : By adjusting  $r$  we can change the “shape” of  $\tilde{J}$  so that it is “close” to  $J^*$  or  $J_\mu$
- A simulator may be used, particularly when there is no mathematical model of the system (but there is a computer model)
- **We will focus on simulation** , but this is not the only possibility
- We may also use parametric approximation for  $Q$ -factors or cost function differences



# Approximation Architectures

Two key issues:

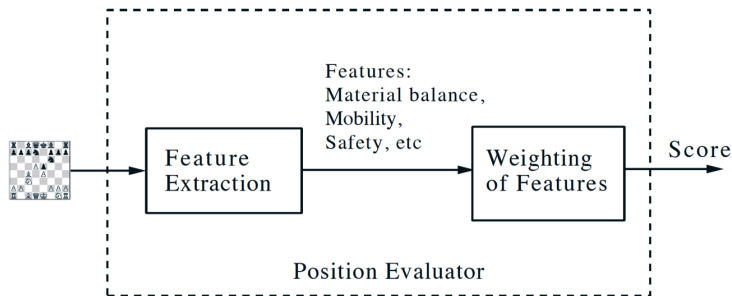
- The choice of parametric class  $\tilde{J}(i; r)$  (the approximation architecture)
- Method for tuning the weights (“training” the architecture)

Success depends strongly on how these issues are handled ... also on insight about the problem

- Divided in linear and nonlinear [i.e., linear or nonlinear dependence of  $\tilde{J}(i; r)$  on  $r$ ]
- Linear architectures are easier to train, but nonlinear ones (e.g., neural networks) are richer

## Computer chess example

- Think of **board position as state** and **move as control**
- Uses a feature-based position evaluator that assigns a score (or approximate  $Q$ -factor) to each position/move



- Relatively few special features and weights, and multistep lookahead

# Linear Approximation Architectures

- With well-chosen features, we can use a linear architecture:

$$\tilde{J}(i; r) = \phi(i)'r, \quad i = 1, \dots, n,$$

or

$$\tilde{J}(r) = \Phi r = \sum_{j=1}^s \Phi_j r_j$$

$\Phi$ : the matrix whose rows are  $\phi(i)'$ ,  $i = 1, \dots, n$ ,  $\Phi_j$  is the  $j$ th column


- This is approximation on the subspace

$$S = \{\Phi r \mid r \in \mathbb{R}^s\}$$

spanned by the columns of  $\Phi$  (basis functions)

# Linear Approximation Architectures

Often, the features encode much of the nonlinearity inherent in the cost function approximated



- Many examples of feature types: Polynomial approximation, radial basis functions, etc

## Example: Polynomial type

- Polynomial Approximation , e.g., a quadratic approximating function. Let the state be  $i = (i_1, \dots, i_q)$  (i.e., have  $q$  “dimensions”) and define

$$\phi_0(i) = 1, \quad \phi_k(i) = i_k, \quad \phi_{km}(i) = i_k i_m, \quad k, m = 1, \dots, q$$

Linear approximation architecture :

$$\tilde{J}(i; r) = r_0 + \sum_{k=1}^q r_k i_k + \sum_{k=1}^q \sum_{m=k}^q r_{km} i_k i_m,$$

where  $r$  has components  $r_0$ ,  $r_k$ , and  $r_{km}$ .


- Interpolation : A subset  $I$  of special/representative states is selected, and the parameter vector  $r$  has one component  $r_i$  per state  $i \in I$ . The approximating function is

$$\tilde{J}(i; r) = r_i, \quad i \in I,$$

$$\tilde{J}(i; r) = \text{interpolation using the values at } i \in I, \quad i \notin I$$

For example, piecewise constant, piecewise linear, more general polynomial interpolations .

## Another Example



- $J^*(i)$ : optimal score starting from position  $i$
- Number of states  $> 2^{200}$  (for  $10 \times 20$  board)
- Success with just 22 features, readily recognized by tetris players as capturing important aspects of the board position (heights of columns, etc)

- 1 Review: Infinite-Horizon DP
- 2 Dimension Reduction in RL
  - Approximation in Value Space
  - Approximation in Policy Space
  - State Aggregation
- 3 On-Policy Learning
  - Direct Projection
  - Bellman Error Minimization
  - Projected Bellman Equation Method
  - From On-Policy to Off-Policy

# Approximation in Policy Space

- A brief discussion; we will return to it later.
- Use parametrization  $\mu(i; r)$  of policies with a vector  $r = (r_1, \dots, r_s)$  .  
Examples:
- Polynomial, e.g.,  $\mu(i; r) = r_1 + r_2 \cdot i + r_3 \cdot i^2$
- Linear feature-based

$$\mu(i; r) = \phi_1(i) \cdot r_1 + \phi_2(i) \cdot r_2$$



# Approximation in Policy Space

- Optimize the cost over  $r$  . For example:
- Each value of  $r$  defines a stationary policy, with cost starting at state  $i$  denoted by  $\tilde{J}(i; r)$ .
- Let  $(p_1, \dots, p_n)$  be some probability distribution over the states, and minimize over  $r$

$$\sum_{i=1}^n p_i \tilde{J}(i; r)$$


- Use a random search, gradient, or other method
- A special case : The parameterization of the policies is indirect, through a cost approximation architecture  $\hat{J}$ , i.e.,

$$\mu(i; r) \in \arg \min_{u \in U(i)} \sum_{j=1}^n p_{ij}(u) (g(i, u, j) + \alpha \hat{J}(j; r))$$

- 1 Review: Infinite-Horizon DP
- 2 Dimension Reduction in RL
  - Approximation in Value Space
  - Approximation in Policy Space
  - State Aggregation
- 3 On-Policy Learning
  - Direct Projection
  - Bellman Error Minimization
  - Projected Bellman Equation Method
  - From On-Policy to Off-Policy


# Aggregation

- A first idea : Group similar states together into “aggregate states”  $x_1, \dots, x_s$ ; assign a common cost value  $r_i$  to each group  $x_i$ .
- Solve an “aggregate” DP problem , involving the aggregate states, to obtain  $r = (r_1, \dots, r_s)$ . This is called hard aggregation



$$\Phi = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

# Aggregation



$$\Phi = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

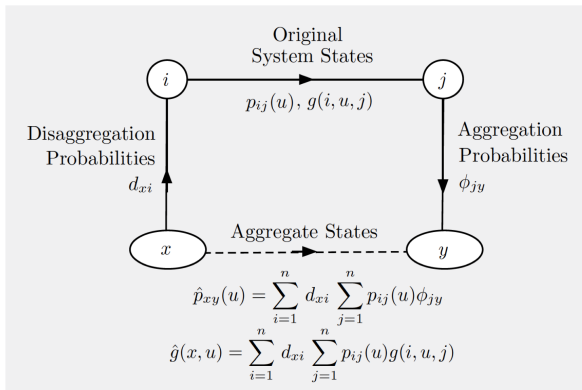
- More general/mathematical view : Solve

$$\Phi r = \Phi D T_{\mu}(\Phi r)$$

where the rows of  $D$  and  $\Phi$  are prob. distributions (e.g.,  $D$  and  $\Phi$  “aggregate” rows and columns of the linear system  $J = T_{\mu}J$ )

- Compare with projected equation  $\Phi r = \Pi T_{\mu}(\Phi r)$ . Note:  $\Phi D$  is a projection in some interesting cases

# Aggregation as Problem Abstraction



- Aggregation can be viewed as a systematic approach for problem approximation. Main elements:
- **Solve (exactly or approximately) the “aggregate” problem** by any kind of VI or PI method (including simulation-based methods)
- **Use the optimal cost of the aggregate problem to approximate** the optimal cost of the original problem

# Aggregate System Description

- The transition probability from aggregate state  $x$  to aggregate state  $y$  under control  $u$

$$\hat{p}_{xy}(u) = \sum_{i=1}^n d_{xi} \sum_{j=1}^n p_{ij}(u) \phi_{jy}, \quad \text{or } \hat{P}(u) = DP(u)\Phi$$

where the rows of  $D$  and  $\Phi$  are the disaggregation and aggregation probs.

- The expected transition cost is

$$\hat{g}(x, u) = \sum_{i=1}^n d_{xi} \sum_{j=1}^n p_{ij}(u) g(i, u, j), \quad \text{or } \hat{g} = DP(u)g$$

# Aggregate Bellman's Equation

- The optimal cost function of the aggregate problem, denoted  $\hat{R}$ , is

$$\hat{R}(x) = \min_{u \in U} \left[ \hat{g}(x, u) + \alpha \sum_y \hat{p}_{xy}(u) \hat{R}(y) \right], \quad \forall x$$

Bellman's equation for the aggregate problem.

- The optimal cost function  $J^*$  of the original problem is approximated by  $\tilde{J}$  given by

$$\tilde{J}(j) = \sum_y \phi_{jy} \hat{R}(y), \quad \forall j$$

## Example I: Hard Aggregation

- Group the original system states into subsets, and view each subset as an aggregate state
- Aggregation probs.:  $\phi_{jy} = 1$  if  $j$  belongs to aggregate state  $y$ .
- Disaggregation probs.: There are many possibilities, e.g., all states  $i$  within aggregate state  $x$  have equal prob.  $d_{xi}$ .
- If optimal cost vector  $J^*$  is piecewise constant over the aggregate states/subsets, hard aggregation is exact. Suggests grouping states with “roughly equal” cost into aggregates.
- A variant: Soft aggregation (provides “soft boundaries” between aggregate states).



## Example II: Feature-Based Aggregation

- Important question: How do we group states together?
- If we know good features, it makes sense to group together states that have “similar features”
- A general approach for passing from a feature-based state representation to a hard aggregation-based architecture
- Essentially discretize the features and generate a corresponding piecewise constant approximation to the optimal cost function
- Aggregation-based architecture is more powerful (it is nonlinear in the features)
- ... but may require many more aggregate states to reach the same level of performance as the corresponding linear feature-based architecture

## Example III: Rep. States/Coarse Grid

- Choose a collection of “representative” original system states, and associate each one of them with an aggregate state
- Disaggregation probabilities are  $d_{xi} = 1$  if  $i$  is equal to representative state  $x$ .
- Aggregation probabilities associate original system states with convex combinations of representative states


$$j \sim \sum_{y \in \mathcal{A}} \phi_{jy} y$$

- Well-suited for Euclidean space discretization
- Extends nicely to continuous state space, including belief space of POMDP

- 1 Review: Infinite-Horizon DP
- 2 Dimension Reduction in RL
  - Approximation in Value Space
  - Approximation in Policy Space
  - State Aggregation
- 3 On-Policy Learning
  - Direct Projection
  - Bellman Error Minimization
  - Projected Bellman Equation Method
  - From On-Policy to Off-Policy

## Direct Policy evaluation

- Approximate the cost of the current policy by using least squares and simulation-generated cost samples
- Amounts to projection of  $J_\mu$  onto the approximation subspace



- Solution by least squares methods
- Regular and optimistic policy iteration
- Nonlinear approximation architectures may also be used

# Direct Evaluation by Simulation

- **Projection by Monte Carlo Simulation:** Compute the projection  $\Pi J_\mu$  of  $J_\mu$  on subspace  $\mathcal{S} = \{\Phi r \mid r \in \mathbb{R}^s\}$ , with respect to a weighted Euclidean norm  $\|\cdot\|_\xi$
- Equivalently, find  $\Phi r^*$ , where

$$r^* = \arg \min_{r \in \mathbb{R}^s} \|\Phi r - J_\mu\|_\xi^2 = \arg \min_{r \in \mathbb{R}^s} \sum_{i=1}^n \xi_i (\phi(i)'r - J_\mu(i))^2$$

- Setting to 0 the gradient at  $r^*$ ,

$$r^* = \left( \sum_{i=1}^n \xi_i \phi(i) \phi(i)' \right)^{-1} \sum_{i=1}^n \xi_i \phi(i) J_\mu(i)$$

## Direct Evaluation by Simulation

- Generate samples  $\{(i_1, J_\mu(i_1)), \dots, (i_k, J_\mu(i_k))\}$  using distribution  $\xi$
- Approximate by Monte Carlo the two “expected values” with low-dimensional calculations

$$\hat{r}_k = \left( \sum_{t=1}^k \phi(i_t) \phi(i_t)' \right)^{-1} \sum_{t=1}^k \phi(i_t) J_\mu(i_t)$$

- Equivalent least squares alternative calculation:

$$\hat{r}_k = \arg \min_{r \in \mathbb{R}^s} \sum_{t=1}^k (\phi(i_t)' r - J_\mu(i_t))^2$$

# Convergence of Evaluated Policy

- By law of large numbers, we have

$$\frac{1}{k} \sum_{t=1}^k \phi(i_t) \phi(i_t)' \xrightarrow{a.s.} \frac{1}{n} \sum_{i=1}^n \xi_i \phi(i) \phi(i)'$$

and

$$\frac{1}{k} \sum_{t=1}^k \phi(i_t) J_{\mu}(i_t) \xrightarrow{a.s.} \frac{1}{n} \sum_{i=1}^n \xi_i \phi(i) J_{\mu}(i)$$

- We have

$$r_k \xrightarrow{a.s.} r^* = \Pi_S J_{\mu}.$$

- As the number of samples increases, the estimated low-dim cost  $r_k$  converges almost surely to the projected  $J_{\mu}$ .

## Indirect policy evaluation

- Solve the **projected equation**

$$\Phi r = \Pi T_\mu(\Phi r)$$

where  $\Pi$  is projection w/ respect to a suitable weighted Euclidean norm

- Solution methods that use simulation (to manage the calculation of  $\Pi$ )
- TD( $\lambda$ ): Stochastic iterative algorithm for solving

$$\Phi r = \Pi T_\mu(\Phi r)$$

- LSTD( $\lambda$ ): Solves a simulation-based approximation w/ a standard solver
- LSPE( $\lambda$ ): A simulation-based form of projected value iteration ; essentially

$$\Phi r_{k+1} = \Pi T_\mu(\Phi r_k) + \text{simulation noise}$$

- Almost sure convergence guarantee



- 1 Review: Infinite-Horizon DP
- 2 Dimension Reduction in RL
  - Approximation in Value Space
  - Approximation in Policy Space
  - State Aggregation
- 3 On-Policy Learning
  - Direct Projection
  - Bellman Error Minimization
  - Projected Bellman Equation Method
  - From On-Policy to Off-Policy

# Bellman equation error methods

- Another example of indirect approximate policy evaluation:

$$\min_r \|\Phi r - T_\mu(\Phi r)\|_\xi^2 \quad (*)$$

where  $\|\cdot\|_\xi$  is Euclidean norm, weighted with respect to some distribution  $\xi$

- It is closely related to the projected equation/Galerkin approach (with a special choice of projection norm)
- Several ways to implement projected equation and Bellman error methods by simulation . They involve:
  - Generating many random samples of states  $i_k$  using the distribution  $\xi$
  - Generating many samples of transitions  $(i_k, j_k)$  using the policy  $\mu$
  - Form a simulation-based approximation of the optimality condition for projection problem or problem (\*) (use sample averages in place of inner products)
  - Solve the Monte-Carlo approximation of the optimality condition
- Issues for indirect methods: How to generate the samples ? How to calculate  $r^*$  efficiently ?

- 1 Review: Infinite-Horizon DP
- 2 Dimension Reduction in RL
  - Approximation in Value Space
  - Approximation in Policy Space
  - State Aggregation
- 3 On-Policy Learning
  - Direct Projection
  - Bellman Error Minimization
  - Projected Bellman Equation Method
  - From On-Policy to Off-Policy

# Cost Function Approximation via Projected Equations

Ideally, we want to solve the Bellman equation (for a fixed policy  $\mu$ )

$$J = T_\mu J$$

In MDP, the equation is  $n \times n$ :

$$J = g_\mu + \alpha P_\mu J.$$

We solve a projected version of the high-dim equation

$$J = \Pi(g_\mu + \alpha P_\mu J)$$

Since the projection  $\Pi$  is onto the space spanned by  $\Phi$ , the projected equation is equivalent to

$$\Phi r = \Pi(g_\mu + \alpha P_\mu \Phi r)$$

We fix the policy  $\mu$  from now on, and omit mentioning it.

# Matrix Form of Projected Equation

- The solution  $\Phi r^*$  satisfies the **orthogonality condition** : The error

$$\Phi r^* - (g + \alpha P \Phi r^*)$$

is “orthogonal” to the subspace spanned by the columns of  $\Phi$ .

- This is written as

$$\Phi' \Xi (\Phi r^* - (g + \alpha P \Phi r^*)) = 0,$$

where  $\Xi$  is the diagonal matrix with the steady-state probabilities  $\xi_1, \dots, \xi_n$  along the diagonal.

- Equivalently,  $C r^* = d$ , where

$$C = \Phi' \Xi (I - \alpha P) \Phi, \quad d = \Phi' \Xi g$$

but **computing  $C$  and  $d$  is HARD** (high-dimensional inner products) .

# Simulation-Based Implementations

- Key idea: Calculate simulation-based approximations based on  $k$  samples

$$C_k \approx C, \quad d_k \approx d$$

- Matrix inversion  $r^* = C^{-1}d$  is approximated by

$$\hat{r}_k = C_k^{-1}d_k$$

This is the **LSTD** (Least Squares Temporal Differences) Method.

- Key fact:  $C_k, d_k$  can be computed with low-dimensional linear algebra (of order  $s$ ; the number of basis functions).

## Simulation Mechanics

- We generate an infinitely long trajectory  $(i_0, i_1, \dots)$  of the Markov chain, so states  $i$  and transitions  $(i, j)$  appear with long-term frequencies  $\xi_i$  and  $p_{ij}$ .
- After generating each transition  $(i_t, i_{t+1})$ , we compute the row  $\phi(i_t)'$  of  $\Phi$  and the cost component  $g(i_t, i_{t+1})$ .
- We form

$$d_k = \frac{1}{k+1} \sum_{t=0}^k \phi(i_t) g(i_t, i_{t+1}) \approx \sum_{i,j} \xi_i p_{ij} \phi(i) g(i, j) = \Phi' \Xi g = d$$

$$C_k = \frac{1}{k+1} \sum_{t=0}^k \phi(i_t) (\phi(i_t) - \alpha \phi(i_{t+1}))' \approx \Phi' \Xi (I - \alpha P) \Phi = C$$

- Convergence based on law of large numbers:  $C_k \xrightarrow{a.s.} C, d_k \xrightarrow{a.s.} d$ . As sample size increases,  $r_k$  converges a.s. to the solution of projected Bellman equation.

- 1 Review: Infinite-Horizon DP
- 2 Dimension Reduction in RL
  - Approximation in Value Space
  - Approximation in Policy Space
  - State Aggregation
- 3 On-Policy Learning
  - Direct Projection
  - Bellman Error Minimization
  - Projected Bellman Equation Method
  - From On-Policy to Off-Policy



# Approx. PI via On-Policy Learning

## Outer Loop (Off-Policy RL):

- Estimate the value function of the current policy  $\mu_t$  using linear features:

$$J_{\mu_t} \approx \Phi r_t.$$

## Inner Loop (On-Policy RL):

- Generate state trajectories ...
- Estimate  $r_t$  via Bellman error minimization (or direct projection, or projected equation approach)
- Update the policy by

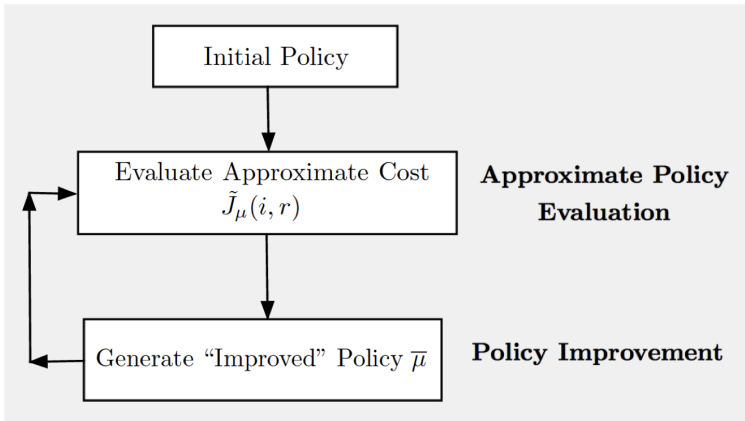
$$\mu_{t+1}(i) = \operatorname{argmin}_a \sum_j p_{ij}(a)(g(i, a, j) + \phi(j)'r_t), \quad \forall i.$$

## Comments:

- Requires knowledge of  $p_{ij}$  (suitable for computer games with known transitions)
- The policy  $\mu_{t+1}$  is parameterized by  $r_t$ .

## Approx. PI via On-Policy Learning

- Use simulation to approximate the cost  $J_\mu$  of the current policy  $\mu$
- Generate “improved” policy  $\bar{\mu}$  by minimizing in (approx.) Bellman equation



Alternatively we can approximate the  $Q$ -factors of  $\mu$

# Theoretical Basis of Approximate PI

- If policies are approximately evaluated using an approximation architecture such that

$$\max_i |\tilde{J}(i, r_k) - J_{\mu^k}(i)| \leq d, \quad k = 0, 1, \dots$$

- If policy improvement is also approximate,

$$\max_i |(T_{\mu^{k+1}} \tilde{J})(i, r_k) - (T \tilde{J})(i, r_k)| \leq \epsilon, \quad k = 0, 1, \dots$$

- **Error bound:** The sequence  $\{\mu^k\}$  generated by approximate policy iteration satisfies

$$\limsup_{k \rightarrow \infty} \max_i (J_{\mu^k}(i) - J^*(i)) \leq \frac{\epsilon + 2\alpha d}{(1 - \alpha)^2}$$

- Typical practical behavior: The method makes steady progress up to a point and then the iterates  $J_{\mu^k}$  oscillate within a neighborhood of  $J^*$ .
- In practice oscillations between policies is probably not the major concern.